

Socket Programming

By: Krishan Kumar Sethi
KIIT, Bhubaneswar

Outline


- Client Server Architecture
- Introduction to Socket
- IP address
- Ports
- Byte Ordering
- Connectionless service
- Connection oriented service
- Various library functions for socket programming

Client/Server Computing

- Some hosts (clients, typically desktop computers) are specialized to interact with users:
 - Gather input from users
 - Present information to users
- Other hosts (servers) are specialized to manage large data, process that data
- The Web is a good example: Client (Browser) & Server (HTTP server)


Client/Server Computing

- Other examples:
 - E-mail




Client/Server Computing

- Other examples:
 - Chatroom




Tiered Client/Server Architecture

- 1-tier: single program
- 2-tier: client/server (e.g. the Web)
- 3-tier: application logic and databases on different servers (e.g. the Web with CGI and databases)



Client/Server Communication

- Two related processes on a single machine may communicate through a pipe



- A pipe is a pseudo-file that can be used to connect two processes together

Client/Server Communication


- Two UNRELATED processes may communicate through files (process A write to a file and process B reads from it)
- But HOW two processes located on two different machines communicate? Solution: Berkeley sockets.

What are sockets

- A socket is an end point of a connection
- Or: the interface between user and network
- Two sockets must be connected before they can be used to transfer data (case of TCP)
- Message destinations are specified as socket addresses.
- Each socket address is a communication identifier:
 - Internet address
 - Port number
- The port number is an integer that is needed to distinguish between services running on the same machine
- A number of connections to choose from:
 - TCP, UDP
- Types of Sockets
 - SOCK_STREAM, SOCK_DGRAM

IP Address

- It is a logical address assigned to each machine to identify it uniquely in the computer network.
- 32-bit identifier
- Dotted-quad: 192.118.56.25
- `www.google.com` -> 167.208.101.28
- Identifies a host interface (not a host)



Ports

- it is used to identify a unique process/service on the server.
- 16 bits unsigned number. Total possible ports are $(0 - (2^{16} - 1)) = (0 - 65535)$
- Port numbers between 0 .. 1023 are reserved for well known applications.
- Some “well-known” ports:
 - 21: ftp
 - 23: telnet
 - 80: http
 - 161: snmp
- Check out `/etc/services` file for complete list of ports and services associated to those ports

Which transport protocol (TCP vs. UDP)

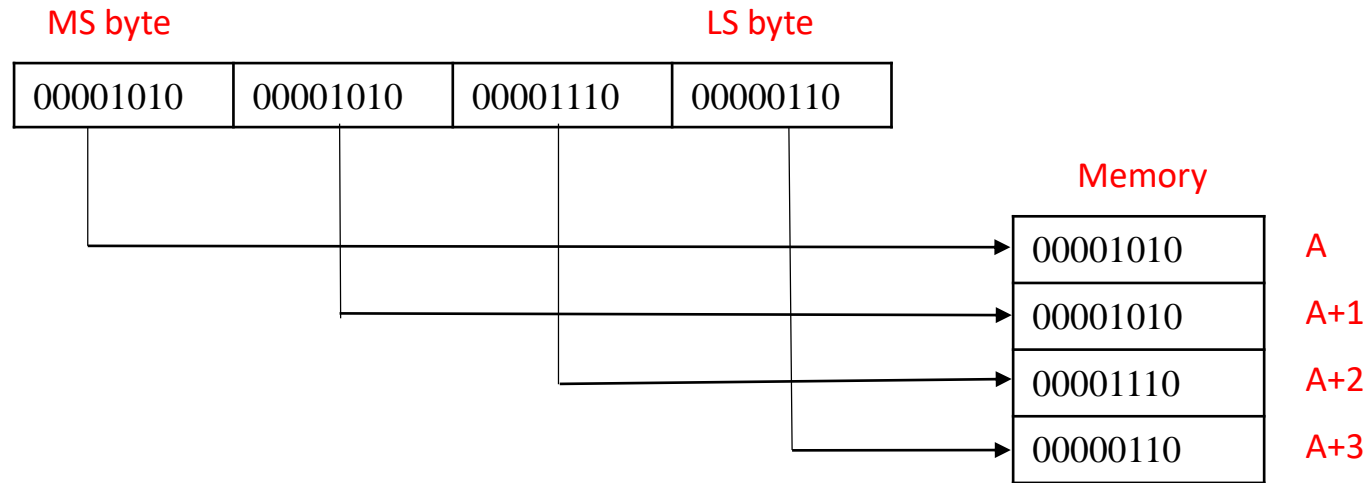
- TCP -- Transmission Control Protocol
- UDP -- User Datagram Protocol
- What should I use?
 - TCP is a ***reliable*** protocol, UDP is not
 - TCP is ***connection-oriented***, UDP is ***connectionless***
 - TCP incurs overheads, UDP incurs fewer overheads
 - UDP has a size limit of 64k, in TCP no limit

Byte Ordering/ Endianness

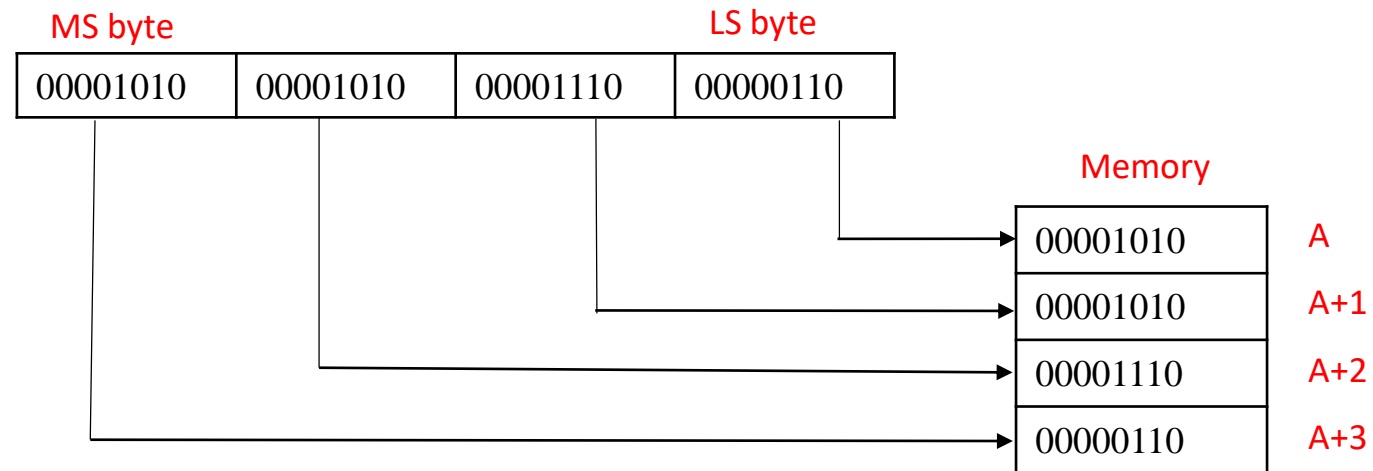
- In computing, endianness is the order or sequence of bytes of a word of digital data in computer memory.
- Types:
 - Big-Endian (Network Byte Order): SUN, PPC MAC, Internet
 - High-order byte of the number is stored in memory at the lowest address.
 - Big End First
 - Little-Endian: x86
 - Low-order byte of the number is stored in memory at the lowest address
 - Low End first
- Network stack (TCP/IP) expects Network Byte Order

Big Endian vs Little Endian

Big Endian:



Little Endian:



Example

- Variable X has 4 byte representation 0x01234567
- Address of X (&X) is given by 0x100
- Big Endian

			0x100	0x101	0x102	0x103			
			01	23	45	67			

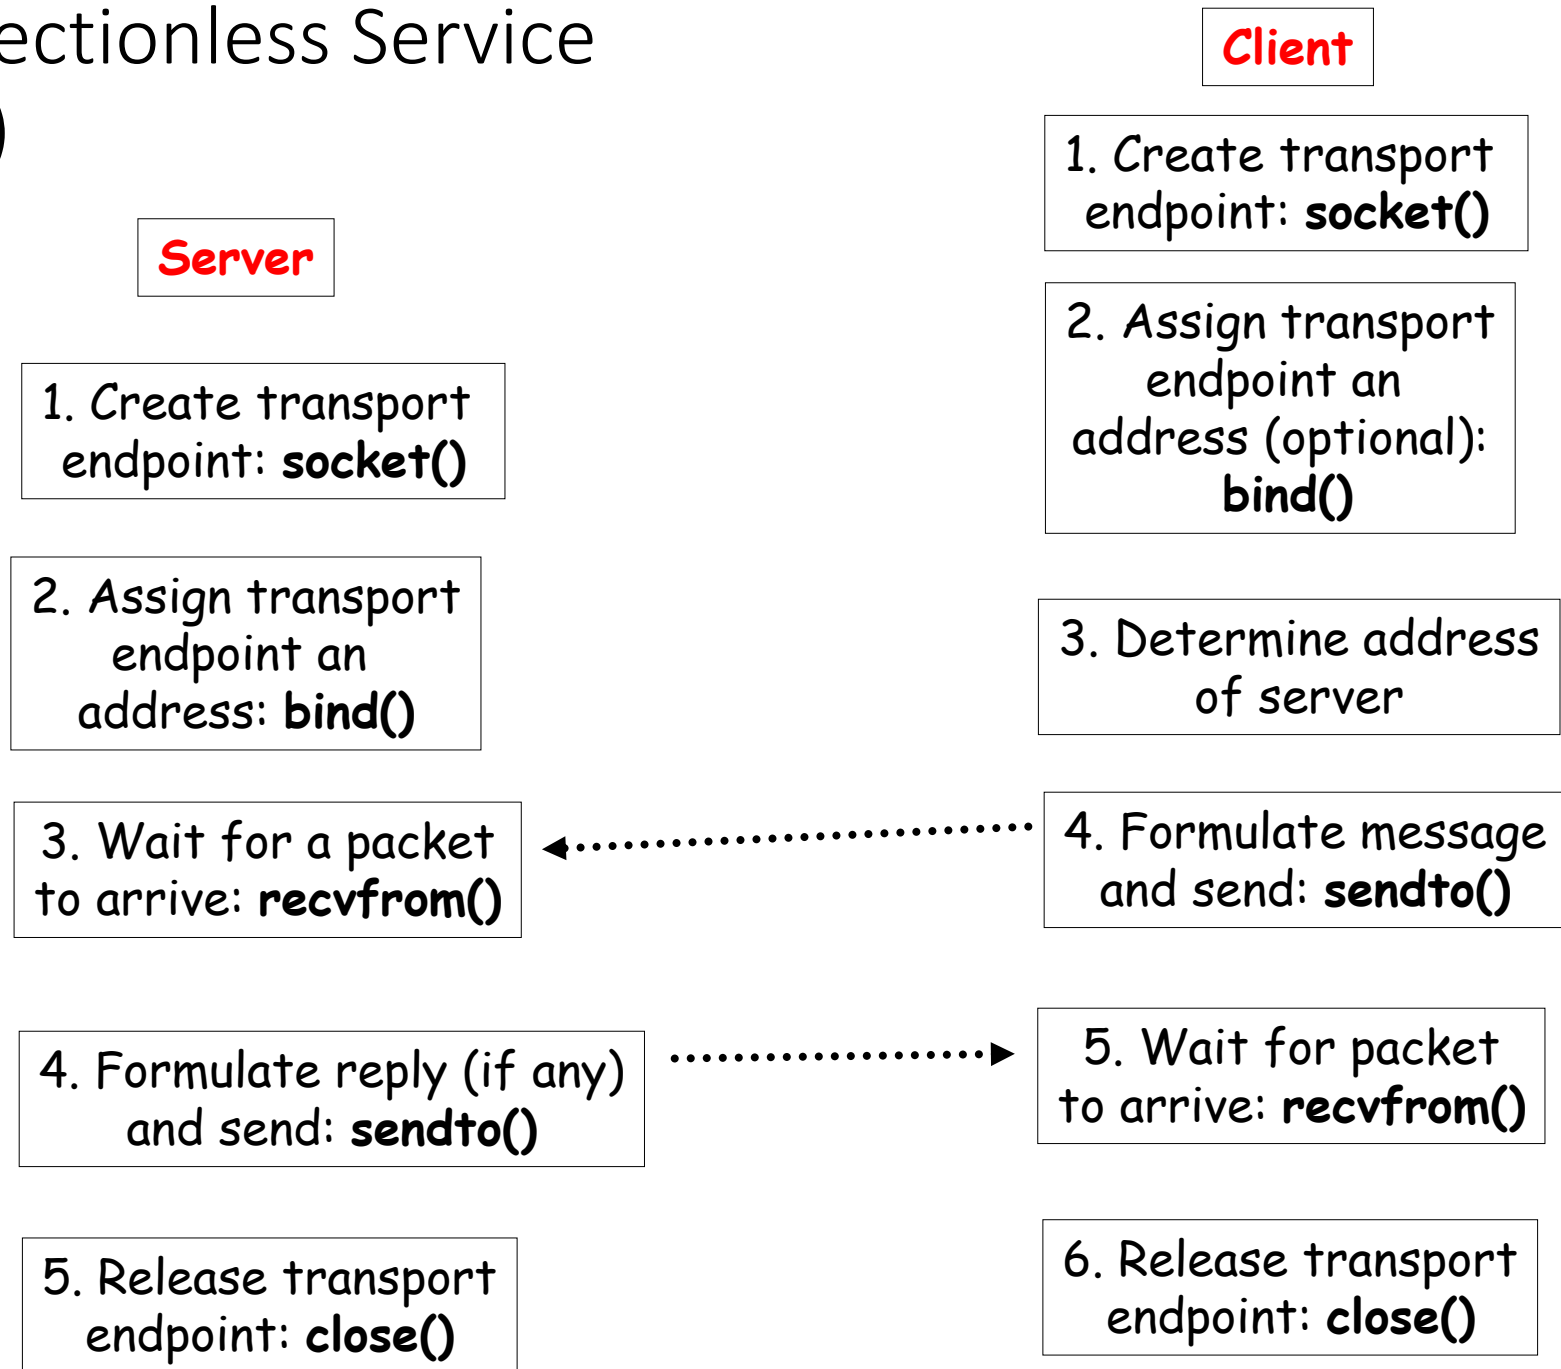
- Little Endian

			0x100	0x101	0x102	0x103			
			67	45	23	01			

Conversions:

- `htonl()`: Translates an unsigned long integer into network byte order.
- `htons()`: Translates an unsigned short integer into network byte order.
- `ntohl()`: Translates an unsigned long integer into host byte order.
- `ntohs()`: Translates an unsigned short integer into host byte order.

Connectionless Service (UDP)



CONNECTION-ORIENTED SERVICE

Server

1. Create transport endpoint for incoming connection request: **socket()**

2. Assign transport endpoint an address: **bind()**

3. Announce willing to accept connections: **listen()**

4. Block and Wait for incoming request: **accept()**

5. Wait for a packet to arrive: **recv ()**

6. Formulate reply (if any) and send: **send()**

7. Release transport endpoint: **close()**

Client

1. Create transport endpoint: **socket()**

2. Assign transport endpoint an address (optional): **bind()**


3. Determine address of server

4. Connect to server: **connect()**

4. Formulate message and send: **send ()**

5. Wait for packet to arrive: **recv()**

6. Release transport endpoint: **close()**



socket()

- Create a socket, giving access to transport layer service.
- Prototype: `int socket(int family, int type, int protocol);`
- *family* is one of
 - AF_INET (IPv4), AF_INET6 (IPv6), AF_LOCAL (local Unix),
 - AF_ROUTE (access to routing tables), AF_KEY (new, for encryption)
- *type* is one of
 - SOCK_STREAM (TCP), SOCK_DGRAM (UDP)
- *protocol* is 0 (used for some raw socket options)
- upon success returns socket descriptor
 - Integer, like file descriptor
 - Return -1 if failure

bind()

- It is used to associate a socket with identifying address (IP + PORT).
- Prototype: `int bind(int sockfd, const struct sockaddr *myaddr, int addrlen);`
- *sockfd* is socket descriptor from `socket()`
- *myaddr* is a pointer to address struct that contains information about IP address and port.
 - if port is 0, then host will pick ephemeral port
- *addrlen* is length of the *myaddr* structure, i.e., `sizeof(struct myaddr)`
- returns 0 if ok, -1 on error
- Re-running a server may result in bind failure
 - Why? Socket still around in kernel using the port
 - Solution: Wait a minute or two or use function `setsockopt()` to clear the socket

Generic

- struct sockaddr
 - { unsigned short sa_family; /* Address family (e.g., AF_INET) */
 - char sa_data[14]; /* Protocol-specific address information */

};

IP Specific

- struct sockaddr_in
 - { unsigned short sin_family; /* Internet protocol (AF_INET) */
 - unsigned short sin_port; /* Port (16-bits) */
 - struct in_addr sin_addr; /* Internet address (32-bits) */
 - char sin_zero[8]; /* Not used */

};

- sin_port and sin_addr must be in **network byte order**.

sockaddr	Family	Blob		
	2 bytes	2 bytes	4 bytes	8 bytes
sockaddr_in	Family	Port	Internet address	Not used

Address Transformation


```
int      inet_aton ( const char  *strptr , struct in_addr *addrptr ) ;
```

```
char     *inet_ntoa (struct in_addr inaddr ) ;
```

Example:

```
struct sockaddr_in my_addr;  
my_addr.sin_family = AF_INET;  
my_addr.sin_port = htons(MYPORT);  
inet_aton("10.0.0.5",&(my_addr.sin_addr));  
memset(&(my_addr.sin_zero),'\0',8);
```

```
// To convert binary IP to string: inet_ntoa()  
printf("%s", inet_ntoa(my_addr.sin_addr));
```



Byte-Manipulation Functions

- In network programming, we often need to initialize a field, copy the contents of one field to another, or compare the contents of two fields.
 - Cannot use string functions (strcpy, strcmp, ...) which assume null character termination.

```
void *memset ( void *dest , int chr , int len ) ;
```

```
void *memcpy ( void *dest , const void *src , int len ) ;
```

```
int memcmp ( const void *first , const void *second , int len ) ;
```

Common ways to IP address assign

1. `my_addr.sin_addr.s_addr = INADDR_ANY; //use my IP adr (127.0.0.1)`
2. `my_addr.sin_addr.s_addr = inet_addr("10.0.0.1"); //assign a specific IP address`

Example:

```
int sockfd;  
struct sockaddr_in my_addr;  
sockfd = socket(PF_INET, SOCK_STREAM, 0);  
my_addr.sin_family = AF_INET; // host byte order  
my_addr.sin_port = htons(MYPORT); // short, network byte order  
my_addr.sin_addr.s_addr = inet_addr("10.0.0.1");  
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct  
bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr));  
/***** Code needs error checking. Don't forget to do that *****/
```


sendto()

- **Prototype:** `int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);`
- `sockfd`: socket descriptor you want to send data to
- `msg` is pointer to the data you want to send
- `len` is the maximum length of the `msg`
- `to` is a pointer to a struct `sockaddr` which contains the destination IP and port
- `tolen` is `sizeof(struct sockaddr)`
- Set flags to zero
- Function returns the number of bytes actually sent or -1 on error

recvfrom()

- **Prototype:** `int recvfrom(int sockfd, void *buf, int len, int flags, struct sockaddr *from, int *fromlen);`
- *sockfd*: socket descriptor to read from
- *buf*: buffer to read the information from
- *len*: maximum length of the buffer
- *flags* set to zero
- *from* is a pointer to a local struct sockaddr that will be filled with IP address and port of the originating machine
- *fromlen* will contain length of address stored in *from*
- *Returns the number of bytes received or -1 on error*

connect()

- Used by connection-oriented clients to connect to server
 - There must already be a socket bound to a connection-oriented service on the fd
 - There must already be a listening socket on the server
 - You pass in the address (IP address, and port number) of the server.
- Used by connectionless clients to specify a “default send to address”
 - Subsequent “writes” or “sends” don’t have to specify a destination address
 - BUT, there really ISN’T any connection established... this is a bad choice of names!
- **Prototype:** `int connect(int sockfd, struct sockaddr *serv_addr, int addrlen)`
- `sockfd` is the socket descriptor returned by `socket()`
- `serv_addr` is pointer to `struct sockaddr` that contains information on destination IP address and port
- `addrlen` is set to `sizeof(struct sockaddr)`
- client doesn’t need `bind()`
 - OS will pick ephemeral port
- returns -1 on error

Example

```
#define DEST_IP "10.2.44.57"
#define DEST_PORT 5000
main(){
    int sockfd;
    struct sockaddr_in dest_addr; // will hold the destination
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
    dest_addr.sin_family = AF_INET; // host byte order
    dest_addr.sin_port = htons(DEST_PORT); // network byte order
    dest_addr.sin_addr.s_addr = inet_addr(DEST_IP);
    memset(&(dest_addr.sin_zero), '\0', 8); // zero the rest of the struct
    connect(sockfd, (struct sockaddr *)&dest_addr, sizeof(struct sockaddr));
    /***** Don't forget error checking *****/
}
```

listen()

- Change socket state for TCP server.

Prototype: `int listen(int sockfd, int backlog);`

- `sockfd` is the socket file descriptor returned by `socket()`
- `backlog` is the number of connections allowed on the incoming queue
- `listen()` returns -1 on error
- Need to call `bind()` before you can `listen()`

accept()

- `accept()` gets the pending connection on the port you are `listen()`ing on.
- **prototype:** `int accept(int sockfd, struct sockaddr cliaddr, socklen_t *addrlen);`
- `sockfd` is the listening socket descriptor
- information about incoming connection is stored in `addr` which is a pointer to a local `struct sockaddr_in`
- `addrlen` is set to `sizeof(struct sockaddr_in)`
- `accept()` returns a new socket file descriptor to use for this accepted connection and `-1` on error
- returns brand new descriptor, created by OS.
- Note: if create new process or thread, can create concurrent server

Example

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define MYPORT 3490 // the port users will be connecting to
#define BACKLOG 10 // pending connections queue will hold

main(){
int sockfd, new_fd; // listen on sock_fd, new connection on new_fd
struct sockaddr_in my_addr; // my address information
struct sockaddr_in their_addr; // connector's address information
int sin_size; sockfd = socket(PF_INET, SOCK_STREAM, 0);
```

Example (*cont..*)

```
my_addr.sin_family = AF_INET; // host byte order
my_addr.sin_port = htons(MYPORT); // short, network byte order
my_addr.sin_addr.s_addr = INADDR_ANY; // auto-fill with my IP
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct
// don't forget your error checking for these calls:
bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr));
listen(sockfd, BACKLOG);
sin_size = sizeof(struct sockaddr_in);
new_fd = accept(sockfd, (struct sockaddr *)&their_addr, &sin_size);
```


send() and recv()

- The two functions are for communicating over stream sockets or connected datagram sockets.

Prototype: `int send(int sockfd, const void *msg, int len, int flags);`

- `sockfd` is the socket descriptor you want to send data to (got from `accept()`)
- `msg` is a pointer to the data you want to send
- `len` is the length of that data in bytes
- set flags to 0 for now
- `send()` returns the number of bytes actually sent (may be less than the number you told it to send) or -1 on error

- Example:

```
char *msg = "hello!";
```

```
int len, bytes_sent;
```

```
.....
```

```
len = strlen(msg);
```

```
bytes_sent = send(sockfd, msg, len 0);
```

send() and recv()

- **Prototype:** `int recv(int sockfd, void *buf, int len, int flags);`
 - `sockfd` is the socket descriptor to read from
 - `buf` is the buffer to read the information into
 - `len` is the maximum length of the buffer
 - set flags to 0 for now
 - `recv()` returns the number of bytes actually read into the buffer or -1 on error
 - If `recv()` returns 0, the remote side has closed connection on you

close()

Prototype: `int close(int sockfd);`

- Closes connection corresponding to the socket descriptor and frees the socket descriptor
- Will prevent any more sends and receives

How to write + compile + run socket programs (Unix)

- Socket programs can be executed either in unix system using GCC compiler or windows using winsock library.
- Write a separate client side program and server side program in any editor program (gedit, vi, nano etc..).
- Compile: `gcc client.c -o client` (on client side)
`gcc server.c -o server` (on server side)
 - Run: `./client` (on client side)
`./server` (on server side)